

# Service management using Docker

# A brief history lesson

- Before virtualization took over, containers were the goto tool for service encapsulation
  - Solaris had Zones
  - FreeBSD and others had jails
  - Linux had LVS
- The motivation was often security, multi-tenancy and resource consolidation

# Container drawbacks

- Complicated to set up
- Limited to the same OS / distribution
- No external API's or surrounding toolchains, local use was assumed
- The paradigm was still: A few big servers that do many, many things

# Then virtualization took over

- With the advent of multi-core CPU's, cheap memory and hardware virtualization support, virtual environments became the norm
- More convenient than containers
  - OS - independent
  - Complete OS stacks
  - Automation features and API's that ultimately led to cloud computing
- A new paradigm: one VM per task

# Cloud Computing and development projects

- Most VMs used in clouds for production are based on robust, long-term releases, such as Ubuntu 12/14.04 or RHEL
- Developers, on the other hand, run Fedora, Linux Mint etc. on their laptops
- It can be challenging to reproduce the same environment in production
  - What works for the developer may not work on a production server

# The shipping problem

- Continuous integration and DevOps requires code to move quickly and easily between environments
- The code itself is in Git, but what about everything else that is needed?
  - Libraries at specific versions
  - Config files
  - Other packages
- The traditional approach would be configuration management, but that is not installed on a developers laptop...

# The developers approach:



# Based on a simple idea:

- The shipping industry had a similar challenge: people tried to ship all kinds of items of various sizes
- Harbor equipment and road/rail transportation was inadequate to adapt to the massive increase in shipping demand
- Solution: Standardize!





# Docker

- A framework for running, managing and sharing of Linux containers
- Built on modern Linux kernel technologies such as cgroups
- A public service Docker Hub provides similar sharing as github
  - Private hub servers are also possible

# The management interface containers were missing?

- Docker has solved the lack of high-level management that was needed to make containers usable for non-experts
- Combined with more secure technologies, it is now a viable alternative to spawning individual VMs for each task

# Docker concepts

- A docker instance is booted from an image
  - The intention is to run only a single command inside the instance
- All changes are stored in an ephemeral filesystem
- Instances are on a local network with port forwarding from the host
- Instances can have access to local folders on the host (called volumes)
- You can commit, push and pull images just like with git

# Docker use-cases

- Rapid testing and sandboxing
  - New instances spawn almost instantaneously
- A way to "ship" code all the way from developers until production
- Lean service encapsulation
- Providing efficient multi-tenancy

# Docker basics

# Docker support

- Docker is available for the many platforms
  - Ubuntu, Debian
  - RedHat, CentOS, Fedora
  - With boot2docker:
    - Mac OS X
    - Windows

# Installing Docker on Ubuntu 14.04

- Docker advocates using either RHEL or Ubuntu for their packages
- Docker is called docker.io on Ubuntu, but lxc-docker is the name of the docker-maintained version
- The simplest way to get the latest  
`curl -sSL https://get.docker.com/ubuntu/ | sudo sh`

# Your first instance

- Check that docker is installed and running:  
`docker ps`
- Launch your first instance:  
`docker run -i -t ubuntu /bin/bash`
- Play around
- In a separate shell, check status once again while the instance is still up:  
`docker ps`



# Things to check out inside the instance

- Do you see an interface?  
try `ifconfig -a` Or `ip a`
- Any running processes?  
try `ps aux`
- Is the instance online?  
try `apt-get update && apt-get install vim`

# What just happened?

- The command downloaded the latest ubuntu image from the Docker Hub
- A new instance was started with an ephemeral filesystem
- The options "-i" and "-t" created an interactive session as well as a pseudo-terminal respectively
- When the shell was closed, the instance finished and 'disappeared'

# Giving the instance a name

- Docker will assign every instance an ID and a random name
- Using the `--name` option, you can provide a name for the instance

```
docker run --name mycontainer -i -t ubuntu /bin/bash
```

- Instance names need to be unique

# Starting a stopped docker instance

- The command `docker ps` only shows running instances
- Use `docker ps -a` to see all instances
- In order to start a stopped instance, run:  
`docker start <name|id>`
- A started instance will re-run the same command as specified previously
  - You may have to press enter once or twice to get the prompt back
  - Note: not all docker images start with a shell

# Daemonizing a container

- If you wish to run a container in the background, add the `-d` option to the run command
- You can check the output inside a container using the command  
`docker logs <id|name>`
  - Adding `-f` will give continuous output, just like `tail -f`
  - Adding `-t` will add timestamps to the output for better debugging

# Inspecting an instance

- The command `docker top <id|name>` will provide a list of processes running inside the instance
- You can also run a task inside a container:  
`docker exec <id|name> command`
- You can either get the output directly by adding `-t -i` or run the command in the background with `-d`
- This is a common way to do maintenance and monitoring on a instance
- You can also run:  
`docker inspect <name|id>`  
`docker inspect -f '{{.NetworkSettings.IPAddress}}' <name|id>`

# Stopping/deleting an instance

- The graceful way to shut down an instance is  
`docker stop <id|name>`
- In order to take it down more dramatically, run:  
`docker kill <id|name>`
- To delete an instance:  
`docker rm <id|name>`

# Automatically restarting instances

- You can specify that an instance should automatically restart if it fails
- Add the `--restart=always` option to the run command
- Alternatively, `--restart=on-failure:3`, which will restart it 3 times



# Docker basic networking

# Networking concepts

- All docker instances are attached to a local, private 172 network
- Default behavior is that no port forwarding is enabled for an instance
- You can specify a port to be exposed when running the instance:
  - -p 80 (expose port 80 on the instance. Will use a host port between 49000 and 49900)
  - -p 80:80 (expose port 80 on the instance. Will use local port 80)

# Example

- Start a docker and expose port 80

```
docker run --name=webserver -t -i -p 80 ubuntu:14.04 /bin/bash
```

- Next we inspect with docker ps

```
root@dockerhost:~# docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
3ce81f679c54	ubuntu:14.04	"/bin/bash"	6 seconds ago	Up 4 seconds	0.0.0.0:49154->80/tcp	webserver

Docker images

# Images

- Images are constructed from filesystem layers on top of a base filesystem (like ubuntu:14.04)
- You can create your own images using two mechanisms
  - commits
  - Docker files (the recommended method)

# Image naming conventions

- Images are usually associated with a docker hub user, a name and a tag
  - Default tag is 'latest'
- For instance:
  - shykes/couchdb
  - ubuntu:12.04
  - fedora:20

# Creating an image with commits

- Start an instance from an existing image:  
`docker run -t -i ubuntu:14.04 /bin/bash`
- Install what you like:  
`apt-get update; apt-get install apache2`
- You can now commit the image:  
`docker commit -m="A webserver image" <id|name>  
webserver`
- Check the existence:  
`docker images`

# Using a docker file

- Docker files are descriptive files that work as recipes for building images
- The convention is that the file is called 'Dockerfile'
- Typically, one creates a folder for the build containing a docker file and all other files you wish to copy into the image
- You then launch the docker build command on the folder



# Build process

- Create a folder:  
`mkdir static_webserver; cd static_webserver`
- Create the docker file Dockerfile:  
# version 0.1  
FROM ubuntu:14.04  
MAINTAINER Kyrre Begnum "kyrre.begnum@hioa.no"  
RUN apt-get update  
RUN apt-get -y install apache2  
RUN echo "Hello world" > /var/www/html/index.html  
EXPOSE 80
- Build the image  
`docker build -t "staticwebserver:v1" .`
- Launch the instance  
`docker run -P -d staticwebserver /usr/sbin/apache2ctl -D FOREGROUND`
- Test: run `docker ps` and find the local port, then do `wget -O - -q http://ip:port`

# The build process in detail

- The build process works in a more layered way than with docker commit
- For every "RUN" command, a new layer is created
- The finished image, therefore, is much smaller than the base image
- Docker will even keep the intermediate images in a cache
- In case one command fails, the process stops and you can run an instance from the intermediate build point and investigate
- If you build the image one more time, the build process is instant

# Docker knows about git

- One exciting feature is that docker can build images based on folders in a git repository
- For example:  
`docker build -t "apachephp:v1" https://git.cs.hioa.no/kyrre/apache\_php\_docker.git`

# Other build features

- CMD - Specify the command to run (can be overridden on the command line)
- WORKDIR - Set the working directory for RUN commands
- ENV - Set environment variables during build and also when creating the instance
- USER - The user to execute the CMD
- VOLUME - Attach a folder to the instance
- ADD - Copy a file into the image (can also be an URL, AND can automatically unpack archives (tar, zip, gzip))
- COPY - Like ADD but with fewer features
- ONBUILD - Triggers that are to be executed during build
- ENTRYPOINT - Overrides the default shell /bin/sh with another binary

# A better webserver example

**FROM** ubuntu:14.04

**MAINTAINER** Kyrre Begnum "[kyrre.begnum@hioa.no](mailto:kyrre.begnum@hioa.no)"

**RUN** apt-get update

**RUN** apt-get install -y apache2

**RUN** apt-get -y install libapache2-mod-php5 php5-mysql

**ENV** APACHE\_RUN\_USER www-data

**ENV** APACHE\_RUN\_GROUP www-data

**ENV** APACHE\_LOG\_DIR /var/log/apache2

**ONBUILD ADD** . /var/www/html/

**EXPOSE** 80

**ENTRYPOINT** ["/usr/sbin/apache2"]

**CMD** [ "-D","FOREGROUND"]

# What if you need to run several processes inside a container?

- Docker does mainly intend to run only a single process per instance
- Services are disabled from starting
- Third-party tools can be used to rather start a suite of processes, like supervisor: <http://supervisord.org>

# Using supervisor to run apache2 and SSH (The Dockerfile)

- These lines install apache2, ssh server and supervisor

```
RUN apt-get update && apt-get install -y openssh-server apache2 supervisor  
RUN mkdir -p /var/lock/apache2 /var/run/apache2 /var/run/sshd /var/log/  
supervisor
```

- Also, copy a supervisor.conf file into the image:

```
COPY supervisord.conf /etc/supervisor/conf.d/supervisord.conf
```

- Expose both ports:

```
EXPOSE 22 80
```

- Run supervisor as the default command

```
CMD [ "/usr/bin/supervisord" ]
```

# supervisord.conf

```
[supervisord]  
nodaemon=true
```

```
[program:sshd]  
command=/usr/sbin/sshd -D
```

```
[program:apache2]  
command=/bin/bash -c "source /etc/apache2/envvars && exec /usr/sbin/apache2 -DFOREGROUND"
```



# Data Volumes

# Volume features

- Volumes are a shared folder between the host and one or more container instances
- You can add volumes with the `-v` option to the `run` command and the `VOLUME` directive in the Dockerfile
- The volumes can be shared between instances and be reused across images

# Example

- Adding a simple volume  
`docker run -i -t -v /opt/data ubuntu /bin/bash`
- Notice, that you now have an extra mountpoint inside the instance with `df -h`
- In order to find the actual location of the folder, use the `docker inspect id` command
- They are usually located under:  
`/var/lib/docker/vfs/dir/`

# Sharing a specific folder

- You can also specify what folder to share to a container  
`docker run -t -i -v /opt/localata:/opt/data ....`
- This can also be done as read-only  
`-v /opt/localdata:/opt/data:ro`
- It also works for files:  
`-v /opt/mydatafile:/opt/data`
- This feature is not available for Dockerfiles, as they are supposed to be more portable

# Volume containers

- One recommended design is to have an instance with one or more volume and then share that volume between other instances
- Create a volume container ('create' will not run it)  
`docker create -v /dbdata --name dbdata training/postgres`
- Create another container with the same volume:  
`docker run -d --volumes-from dbdata --name db1 training/postgres`
- You can have multiple instances share the same volume that way

# Discussion

# What have we learned?

- Docker is a simple framework for running, managing and building Linux containers
- Makes some tasks really easy
- Can significantly reduce your cloud VM footprint
- Docker is meant to be a piece of a workflow

# Docker rocks the boat

- It's no secret, Docker challenges our existing paradigm of "One VM per task"
- Where do all our operations support tools go?
  - Configuration management
  - Monitoring
  - Backup
  - Logging



# Is docker a new config system?

- A Dockerfile specifies many of the same things a puppet manifest would, is it a replacement?
- Can one be managed by the other?
  - Puppet apply a viable option again?
- The discussion is still ongoing...

# Docker is still very single-host

- Hard to envision a distributed docker yet
- Docker Hub and private registries make distributing images easy, but what about coordination?
- Docker seems to be working on that: Docker swarms and Docker compose

# Docker user interfaces

- Shipyard: <https://github.com/ehazlett/shipyard>
- DockerUI
- maDocker